

Unraveling Misinformation: LSTM-based Classification of Political Statements

Aarush Kartik

Insprit AI

October 20, 2024

Outline

Introduction

Background

Problem Statement

Methodology

Data

Model

Results

Conclusion

Acknowledgments

Introduction

- ▶ The rapid spread of misinformation during events such as presidential elections has intensified the need for reliable fake news detection methods.
- ▶ Misinformation about political candidates, their background, and affiliations tends to spread across social media and news outlets.
- ▶ Such misinformation complicates decision-making processes for individuals and groups, potentially undermining democratic processes.
- ▶ Machine learning techniques have emerged as powerful tools to tackle the challenge of identifying and mitigating fake news.

Literature Review

- ▶ Key works in your field
- ▶ Gaps in existing research
- ▶ Positioning of your work in relation to existing literature

Problem Statement

- ▶ Define the problem you are addressing
- ▶ Why this problem is important
- ▶ Key challenges involved

- ▶ Approach taken to solve the problem
- ▶ Data collection and experimental setup
- ▶ Models, algorithms, or frameworks used
- ▶ Steps involved in your research

Dataset Description

▶ **LIAR Dataset** [?]

- ▶ 12.8K manually labeled short statements
- ▶ Sources include political debates, TV ads, social media posts
- ▶ Each statement annotated with one of six truthfulness ratings:
 - ▶ true
 - ▶ mostly-true
 - ▶ half-true
 - ▶ barely-true
 - ▶ false
 - ▶ pants-fire
- ▶ Includes metadata:
 - ▶ Subject
 - ▶ Speaker
 - ▶ Job Title
 - ▶ State
 - ▶ Party Affiliation

Data Acquisition

- ▶ Obtained from the official repository [?]
- ▶ Developed an automated script for:
 - ▶ Downloading and extracting the dataset
 - ▶ Ensuring reproducibility and efficiency
 - ▶ Checking for local existence to avoid redundant downloads
- ▶ Dataset format:
 - ▶ Tab-separated values (TSV)
 - ▶ Split into training, validation, and test sets

Data Preprocessing

- ▶ Performed several preprocessing steps:
 - ▶ Data loading and column renaming
 - ▶ Statement length computation
 - ▶ Label encoding
 - ▶ Textual data transformation
 - ▶ Data reshaping for LSTM input
 - ▶ One-hot encoding of labels

Data Loading and Column Renaming

- ▶ Used pandas library to read TSV files into DataFrames
- ▶ Renamed columns for clarity:
 - ▶ **ID**: Unique identifier
 - ▶ **Label**: Truthfulness rating
 - ▶ **Statement**: Text content
 - ▶ **Subject, Speaker**
 - ▶ **Job Title, State Info**
 - ▶ **Party Affiliation**

Statement Length Computation

- ▶ Calculated length of each statement
- ▶ Purpose:
 - ▶ Understand distribution of statement lengths
 - ▶ Assess complexity and variability of textual data
- ▶ Computation:

$$l_i = \text{len}(s_i) \quad (1)$$

- ▶ Where:
 - ▶ l_i : Length of the i -th statement
 - ▶ s_i : Text content of the i -th statement

Label Encoding

- ▶ Converted categorical labels to numerical values
- ▶ Mapping:

 true \rightarrow 0
mostly-true \rightarrow 1
 half-true \rightarrow 2
barely-true \rightarrow 3
 false \rightarrow 4
pants-fire \rightarrow 5

- ▶ Reflects ordinal relationship among labels
- ▶ Encoded label for i -th statement: y_i

Textual Data Transformation

- ▶ Employed Bag-of-Words (BoW) model using Count Vectorization
- ▶ Transformed each statement s_i into vector $\mathbf{x}_i \in \mathbb{R}^V$
- ▶ V : Size of vocabulary from training data
- ▶ Components:

$$x_{ij} = \text{Count}(\text{word}_j, s_i) \quad (2)$$

- ▶ Where x_{ij} : Frequency of j -th word in i -th statement

Data Reshaping for LSTM Input

- ▶ LSTM expects input in 3D format:
(samples, timesteps, features)
- ▶ Reshaped feature vectors accordingly:
 - ▶ Number of timesteps set to 1
 - ▶ Each statement treated as a single timestep with V features

One-Hot Encoding of Labels

- ▶ Converted numerical labels y_i into one-hot encoded vectors $\mathbf{y}_i \in \mathbb{R}^6$
- ▶ Representation:

$$\mathbf{y}_i = [y_{i0}, y_{i1}, y_{i2}, y_{i3}, y_{i4}, y_{i5}] \quad (3)$$

- ▶ Where:

$$y_{ij} = \begin{cases} 1 & \text{if } j = y_i \\ 0 & \text{otherwise} \end{cases} \quad (4)$$

- ▶ Suitable for multi-class classification with categorical cross-entropy loss

Data Splitting

- ▶ Dataset already partitioned by authors
- ▶ Splits:
 - ▶ **Training Set:** Train model parameters
 - ▶ **Validation Set:** Hyperparameter tuning, prevent overfitting
 - ▶ **Test Set:** Assess final performance

Summary of Data Preparation

1. **Data Acquisition:** Automated download and extraction
2. **Data Loading:** Read TSV files into DataFrames
3. **Data Cleaning:** Renamed columns, ensured data integrity
4. **Feature Engineering:**
 - ▶ Computed statement lengths
 - ▶ Transformed text into numerical features using BoW
5. **Label Encoding:** Converted labels to numerical and one-hot formats
6. **Data Reshaping:** Adjusted data shape for LSTM input
 - ▶ Prepared data optimally for training the LSTM network
 - ▶ Facilitated efficient learning and improved potential for accurate classification

Model Selection and Description

- ▶ Employed a **Long Short-Term Memory (LSTM)** network for fake news detection.
- ▶ Dataset: *Liar, Liar Pants on Fire* [?].
- ▶ LSTM networks are a type of Recurrent Neural Network (RNN) capable of learning long-term dependencies.
- ▶ Suitable for processing natural language text due to ability to capture sequential patterns.

Model Architecture

- ▶ Designed to effectively capture sequential patterns in textual data.
- ▶ Consists of the following layers:
 - 1. First LSTM Layer**
 - ▶ 50 units
 - ▶ `return_sequences=True`
 - ▶ Returns hidden state output for each input time step.
 - 2. First Dropout Layer**
 - ▶ Dropout rate of 0.7
 - ▶ Prevents overfitting by randomly deactivating 70% of neurons during training.
 - 3. Second LSTM Layer**
 - ▶ 50 units
 - ▶ Processes sequence output from previous layer.
 - 4. Second Dropout Layer**
 - ▶ Dropout rate of 0.7
 - 5. Dense Output Layer**
 - ▶ Fully connected layer with 6 units (corresponding to 6 classes).
 - ▶ Uses softmax activation function.

Model Compilation and Training

- ▶ Compiled with:
 - ▶ **Loss Function:** Categorical Cross-Entropy
 - ▶ **Optimizer:** Adam optimizer [?]
 - ▶ **Learning Rate:** 0.001
- ▶ Trained for up to 30 epochs.
- ▶ Implemented **Early Stopping** based on validation loss to prevent overfitting.

Mathematical Formulation (LSTM Equations)

LSTM Units are defined by the following equations at each time step t :

$$\text{Input Gate: } i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i) \quad (5)$$

$$\text{Forget Gate: } f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f) \quad (6)$$

$$\text{Cell Candidate: } \tilde{C}_t = \tanh(W_c x_t + U_c h_{t-1} + b_c) \quad (7)$$

$$\text{Cell State Update: } C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t \quad (8)$$

Mathematical Formulation (Continued)

$$\text{Output Gate: } o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o) \quad (9)$$

$$\text{Hidden State: } h_t = o_t \odot \tanh(C_t) \quad (10)$$

Where:

- ▶ x_t : Input vector at time t
- ▶ h_{t-1} : Hidden state from previous time step
- ▶ i_t, f_t, o_t : Input, forget, and output gates
- ▶ C_t : Cell state at time t
- ▶ \tilde{C}_t : Candidate cell state
- ▶ σ : Sigmoid activation function
- ▶ \tanh : Hyperbolic tangent function
- ▶ \odot : Element-wise multiplication
- ▶ W, U : Weight matrices
- ▶ b : Bias vectors

Hyperparameters and Regularization

- ▶ **Number of LSTM Units:** 50 units in each LSTM layer
- ▶ **Dropout Rate:** 0.7 after each LSTM layer
- ▶ **Optimizer:** Adam optimizer
- ▶ **Learning Rate:** 0.001
- ▶ **Loss Function:** Categorical Cross-Entropy
- ▶ **Early Stopping:**
 - ▶ Training halted if validation loss does not improve for 5 consecutive epochs
 - ▶ Helps prevent overfitting

Model Training and Evaluation

▶ **Training Procedure:**

- ▶ Mini-batch gradient descent
- ▶ Batch size: 32
- ▶ Monitored performance on validation set
- ▶ Restored best model weights using early stopping

▶ **Evaluation Metrics:**

- ▶ **Accuracy:** Overall correctness
- ▶ **Confusion Matrix:** Detailed breakdown per class
- ▶ **Classification Report:**
 - ▶ Precision
 - ▶ Recall
 - ▶ F1-score

Rationale for Model Selection

- ▶ LSTMs are effective for modeling sequential data and capturing long-term dependencies.
- ▶ Crucial for understanding context and nuances in natural language.
- ▶ Dropout layers serve as regularization to prevent overfitting.
- ▶ Aimed to generalize well to unseen data by learning robust features.

Implementation Details

- ▶ Implemented using **TensorFlow Keras API**.
- ▶ **Model Definition:**
 - ▶ Used `Sequential` API to stack layers linearly.
- ▶ **Callbacks:**
 - ▶ Early stopping implemented using `EarlyStopping` callback.
 - ▶ Monitored validation loss.
- ▶ **Evaluation Functions:**
 - ▶ Custom methods defined for evaluation.
 - ▶ `evaluate` method computes accuracy, confusion matrix, classification report.

- ▶ Evaluated the LSTM model against several baseline models.
- ▶ Utilized standard metrics:
 - ▶ Accuracy
 - ▶ Precision
 - ▶ Recall
 - ▶ F1-score
- ▶ Marginal improvements in accuracy observed.
- ▶ Precision and recall values were significantly lower than anticipated.

LSTM Model Performance

- ▶ **Precision:** 0.21
- ▶ **Recall:** 0.20
- ▶ Significant challenges in identifying true positive instances of fake news.
- ▶ Model struggled to minimize false negatives.
- ▶ Low precision suggests legitimate news articles were misclassified as fake.

Confusion Matrix

	true	mostly-true	half-true	barely-true	false	pants-fire
true	62	37	41	17	41	10
mostly-true	67	46	48	34	38	8
half-true	48	53	60	47	39	18
barely-true	30	28	46	46	45	17
false	45	32	46	36	74	16
pants-fire	13	9	20	15	17	18

Table: Confusion Matrix: True Labels vs. Predicted Labels

Analysis of Results

- ▶ Performance underscores the complexity of fake news detection.
- ▶ Highlights the need for more nuanced feature extraction.
- ▶ Suggests re-evaluation of the training data.
- ▶ Indicates that relying solely on textual features may not suffice.

Conclusion

- ▶ Explored the effectiveness of an LSTM model for fake news detection.
- ▶ Model showed marginal improvement in accuracy over traditional methods.
- ▶ Overall performance was disappointing, particularly in precision and recall.
- ▶ Results reveal complexities inherent in this challenging task.

Future Work

- ▶ Need for a more robust approach to fake news detection.
- ▶ Consider incorporating various data sources and model architectures.
- ▶ Experiment with hybrid models combining LSTMs with:
 - ▶ Convolutional Neural Networks (CNNs)
 - ▶ Graph-based models
- ▶ Aim to capture both sequential and relational features inherent in news articles and their dissemination networks.

Acknowledgments

- ▶ Thanks to your advisors, collaborators, and mentors
- ▶ Acknowledge any funding or institutional support

Thank you!

Questions?